

house-prices-advanced-regression

December 9, 2021

1 Imports / Read Data

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline

from pprint import pprint

from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Lasso
from sklearn.feature_selection import SelectFromModel
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, mean_absolute_error, \
    ↪ explained_variance_score

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Activation, Dropout
from tensorflow.keras.optimizers import Adam

pd.pandas.set_option('display.max_columns', None)
```

```
[2]: train = pd.read_csv('train.csv')
test = pd.read_csv('test.csv')
```

2 EDA

```
[3]: numerical_cont_features = [col for col in train.drop(['Id', 'SalePrice'], \
    ↪ axis=1) if
                                     train[col].dtype in ['int64', 'float64'] and train[col].
    ↪ unique() > 15]
```

```

numerical_cat_features = [col for col in train.drop(['Id', 'SalePrice'],
→axis=1) if
                        train[col].dtype in ['int64', 'float64'] and train[col].
→nunique() <= 15]
categorical_features = [col for col in train.drop('SalePrice', axis=1) if
                        train[col].dtype in ['object']]
numerical_cols = numerical_cont_features + numerical_cat_features

```

```

[4]: null_cols_train = [[col, train[col].isnull().sum(), \
                        round(train[col].isnull().sum() / len(train[col]), 2)] for
→col in \
                        train if train[col].isnull().sum() > 0]
pprint(null_cols_train)
print(f'null_cols_train length: {len(null_cols_train)}')

```

```

[['LotFrontage', 259, 0.18],
 ['Alley', 1369, 0.94],
 ['MasVnrType', 8, 0.01],
 ['MasVnrArea', 8, 0.01],
 ['BsmtQual', 37, 0.03],
 ['BsmtCond', 37, 0.03],
 ['BsmtExposure', 38, 0.03],
 ['BsmtFinType1', 37, 0.03],
 ['BsmtFinType2', 38, 0.03],
 ['Electrical', 1, 0.0],
 ['FireplaceQu', 690, 0.47],
 ['GarageType', 81, 0.06],
 ['GarageYrBlt', 81, 0.06],
 ['GarageFinish', 81, 0.06],
 ['GarageQual', 81, 0.06],
 ['GarageCond', 81, 0.06],
 ['PoolQC', 1453, 1.0],
 ['Fence', 1179, 0.81],
 ['MiscFeature', 1406, 0.96]]
null_cols_train length: 19

```

```

[5]: null_cols_test = [[col, test[col].isnull().sum(), \
                        round(test[col].isnull().sum() / len(test[col]), 2)] for
→col in \
                        test if test[col].isnull().sum() > 0]
pprint(null_cols_test)
print(f'null_cols_test length: {len(null_cols_test)}')

```

```

[['MSZoning', 4, 0.0],
 ['LotFrontage', 227, 0.16],
 ['Alley', 1352, 0.93],
 ['Utilities', 2, 0.0],

```

```

['Exterior1st', 1, 0.0],
['Exterior2nd', 1, 0.0],
['MasVnrType', 16, 0.01],
['MasVnrArea', 15, 0.01],
['BsmtQual', 44, 0.03],
['BsmtCond', 45, 0.03],
['BsmtExposure', 44, 0.03],
['BsmtFinType1', 42, 0.03],
['BsmtFinSF1', 1, 0.0],
['BsmtFinType2', 42, 0.03],
['BsmtFinSF2', 1, 0.0],
['BsmtUnfSF', 1, 0.0],
['TotalBsmtSF', 1, 0.0],
['BsmtFullBath', 2, 0.0],
['BsmtHalfBath', 2, 0.0],
['KitchenQual', 1, 0.0],
['Functional', 2, 0.0],
['FireplaceQu', 730, 0.5],
['GarageType', 76, 0.05],
['GarageYrBlt', 78, 0.05],
['GarageFinish', 78, 0.05],
['GarageCars', 1, 0.0],
['GarageArea', 1, 0.0],
['GarageQual', 78, 0.05],
['GarageCond', 78, 0.05],
['PoolQC', 1456, 1.0],
['Fence', 1169, 0.8],
['MiscFeature', 1408, 0.97],
['SaleType', 1, 0.0]]
null_cols_test length: 33

```

```

[6]: null_cat_cols_train = [[col, train[col].isnull().sum(), \
                             round(train[col].isnull().sum() / len(train[col]), 2)]
    ↪for \
        col in train[categorical_features] if train[col].
    ↪isnull().sum() > 0]
pprint(null_cat_cols_train)
print(f'null_cat_cols_train length: {len(null_cat_cols_train)}')

```

```

[['Alley', 1369, 0.94],
 ['MasVnrType', 8, 0.01],
 ['BsmtQual', 37, 0.03],
 ['BsmtCond', 37, 0.03],
 ['BsmtExposure', 38, 0.03],
 ['BsmtFinType1', 37, 0.03],
 ['BsmtFinType2', 38, 0.03],
 ['Electrical', 1, 0.0],
 ['FireplaceQu', 690, 0.47],

```

```

['GarageType', 81, 0.06],
['GarageFinish', 81, 0.06],
['GarageQual', 81, 0.06],
['GarageCond', 81, 0.06],
['PoolQC', 1453, 1.0],
['Fence', 1179, 0.81],
['MiscFeature', 1406, 0.96]]
null_cat_cols_train length: 16

```

```

[7]: null_cat_cols_test = [[col, test[col].isnull().sum(), \
                           round(test[col].isnull().sum() / len(test[col]), 2)]_
    ↪for \
                           col in test[categorical_features] if test[col].isnull().
    ↪sum() > 0]
pprint(null_cat_cols_test)
print(f'null_cat_cols_test length: {len(null_cat_cols_test)}')

```

```

[['MSZoning', 4, 0.0],
['Alley', 1352, 0.93],
['Utilities', 2, 0.0],
['Exterior1st', 1, 0.0],
['Exterior2nd', 1, 0.0],
['MasVnrType', 16, 0.01],
['BsmtQual', 44, 0.03],
['BsmtCond', 45, 0.03],
['BsmtExposure', 44, 0.03],
['BsmtFinType1', 42, 0.03],
['BsmtFinType2', 42, 0.03],
['KitchenQual', 1, 0.0],
['Functional', 2, 0.0],
['FireplaceQu', 730, 0.5],
['GarageType', 76, 0.05],
['GarageFinish', 78, 0.05],
['GarageQual', 78, 0.05],
['GarageCond', 78, 0.05],
['PoolQC', 1456, 1.0],
['Fence', 1169, 0.8],
['MiscFeature', 1408, 0.97],
['SaleType', 1, 0.0]]
null_cat_cols_test length: 22

```

```

[8]: null_num_cols_train = [[col, train[col].isnull().sum(), \
                             round(train[col].isnull().sum() / len(train[col]), 2)]_
    ↪for \
                             col in train[numerical_cols] if train[col].isnull().
    ↪sum() > 0]
pprint(null_num_cols_train)

```

```
print(f'null_num_cols_train length: {len(null_num_cols_train)}')
```

```
[['LotFrontage', 259, 0.18], ['MasVnrArea', 8, 0.01], ['GarageYrBlt', 81, 0.06]]
null_num_cols_train length: 3
```


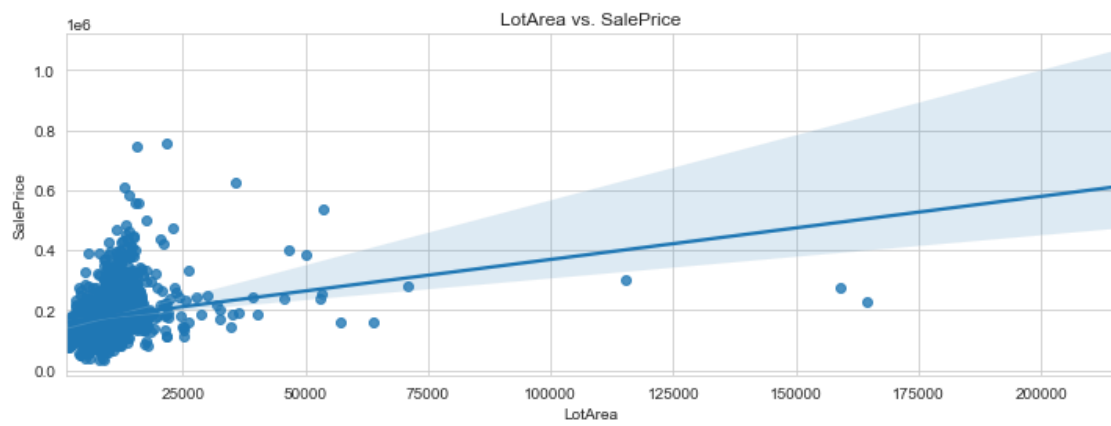

```
[9]: null_num_cols_test = [[col, test[col].isnull().sum(), \
                           round(test[col].isnull().sum() / len(test[col]), 2)]
    ↪ for \
        col in test[numerical_cols] if test[col].isnull().sum()
    ↪ > 0]
pprint(null_num_cols_test)
print(f'null_num_cols_test length: {len(null_num_cols_test)}')
```



```
[['LotFrontage', 227, 0.16],
 ['MasVnrArea', 15, 0.01],
 ['BsmtFinSF1', 1, 0.0],
 ['BsmtFinSF2', 1, 0.0],
 ['BsmtUnfSF', 1, 0.0],
 ['TotalBsmtSF', 1, 0.0],
 ['GarageYrBlt', 78, 0.05],
 ['GarageArea', 1, 0.0],
 ['BsmtFullBath', 2, 0.0],
 ['BsmtHalfBath', 2, 0.0],
 ['GarageCars', 1, 0.0]]
null_num_cols_test length: 11
```

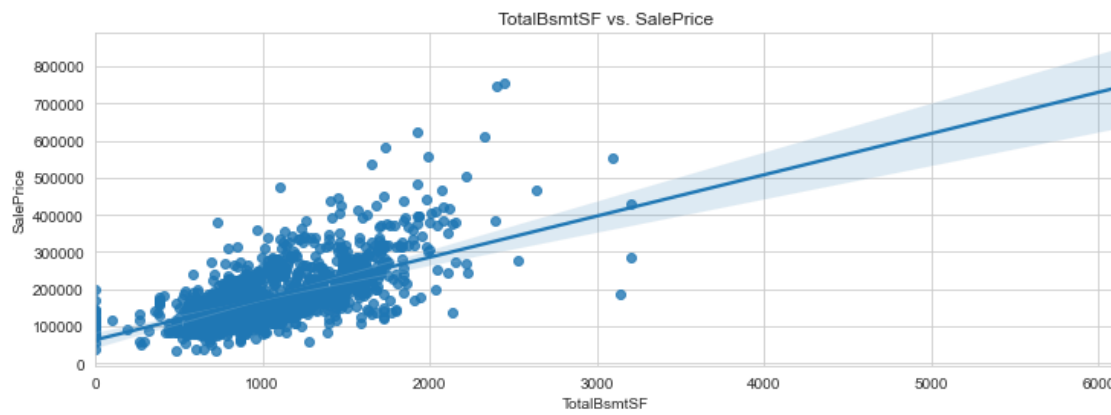
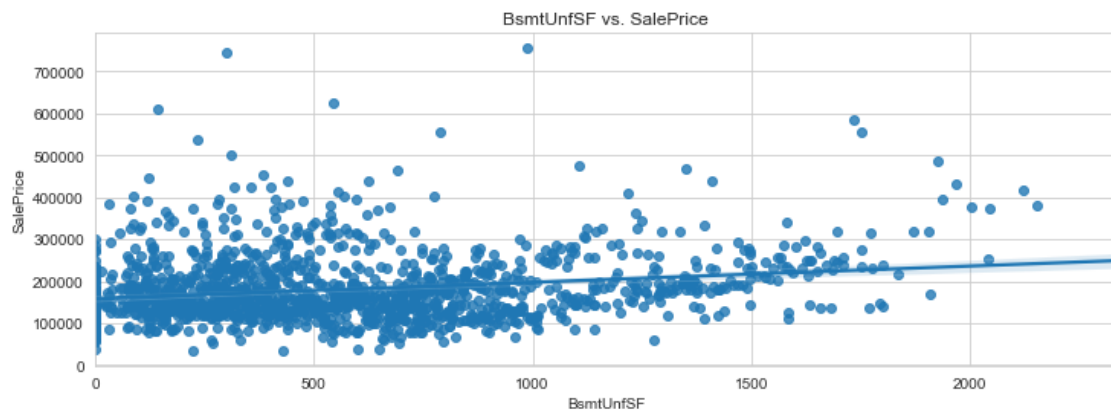
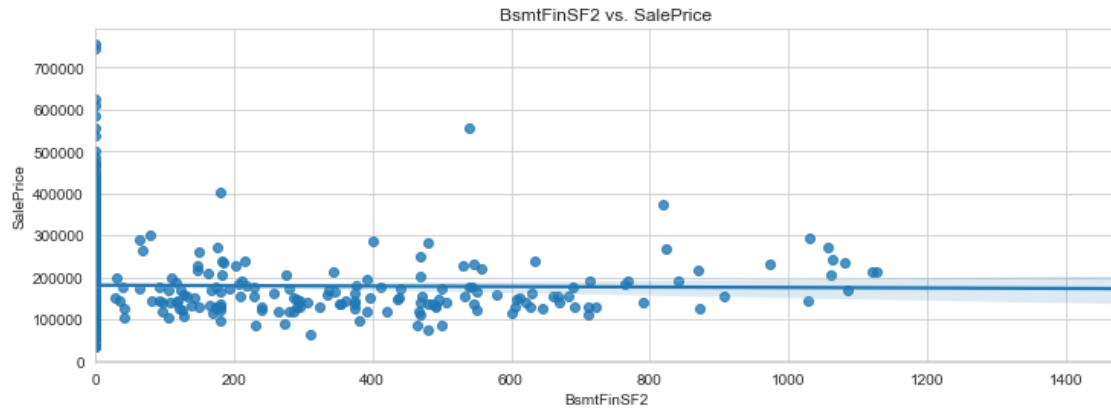
```
[10]: sns.set_style('whitegrid')
```




```
[11]: for ax in train[numerical_cont_features]:
    plt.figure(figsize=(12, 4))
    sns.regplot(x=train[ax], y=train['SalePrice'])
    plt.title(f'{ax} vs. SalePrice')
    plt.xlabel(ax)
    plt.ylabel('SalePrice')
```




```
<ipython-input-11-ed337a0eae73>:2: RuntimeWarning: More than 20 figures have
been opened. Figures created through the pyplot interface
(`matplotlib.pyplot.figure`) are retained until explicitly closed and may
consume too much memory. (To control this warning, see the rcParam
`figure.max_open_warning`).
    plt.figure(figsize=(12, 4))
```

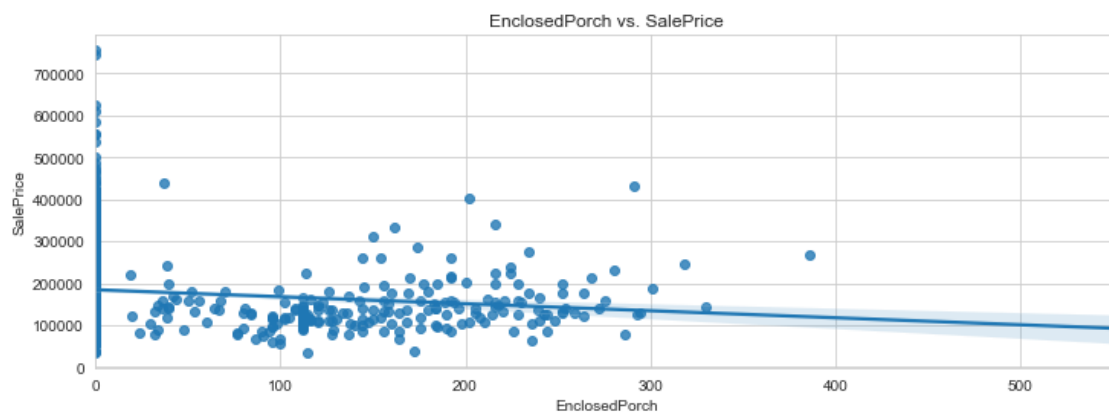

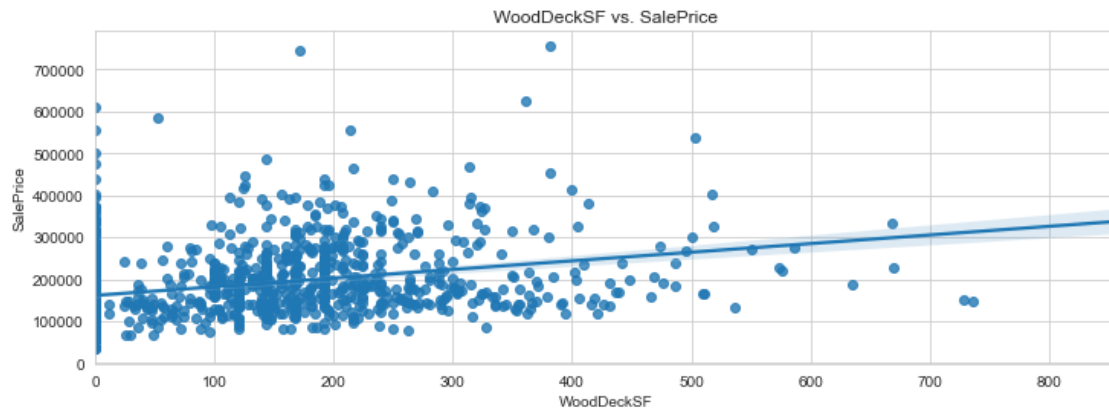


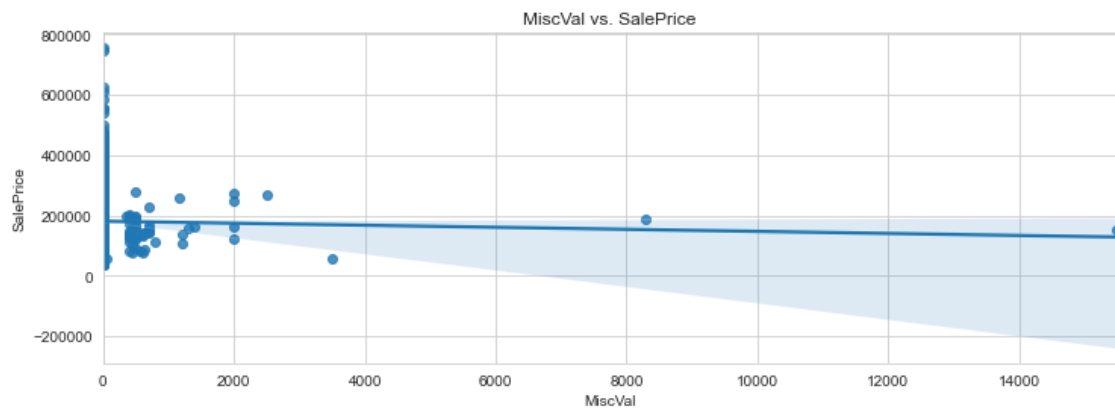

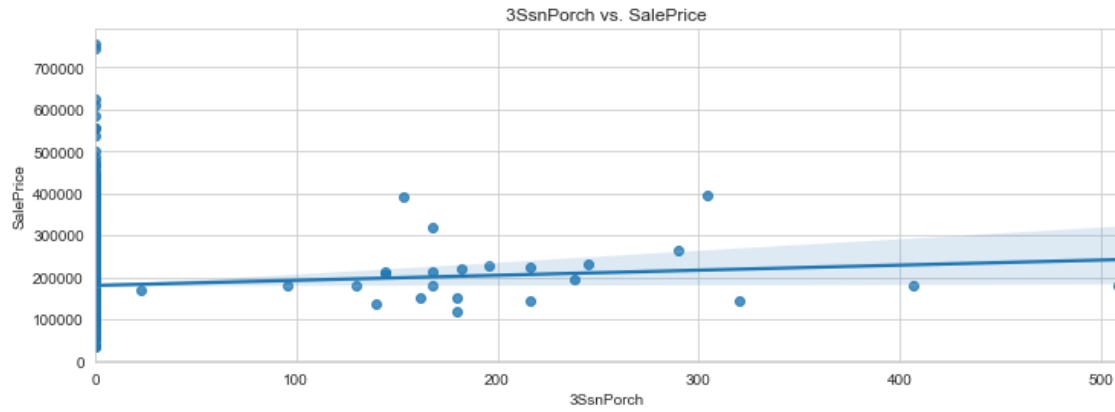






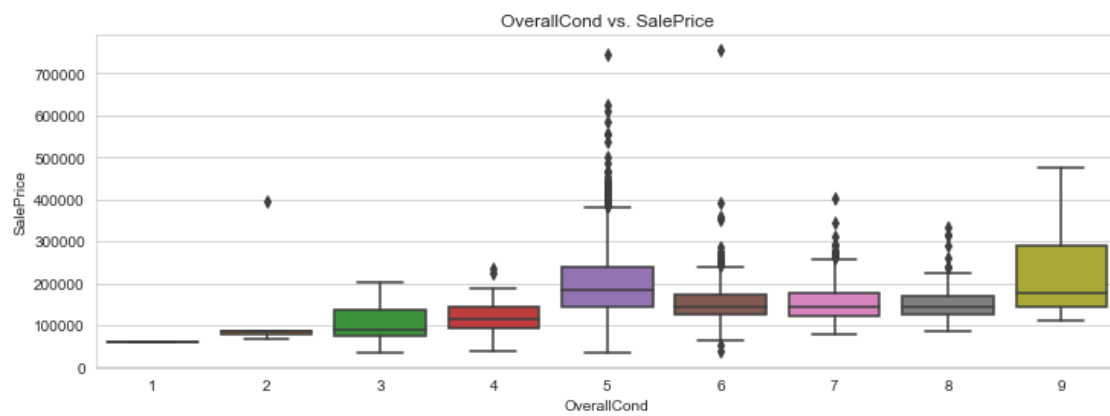

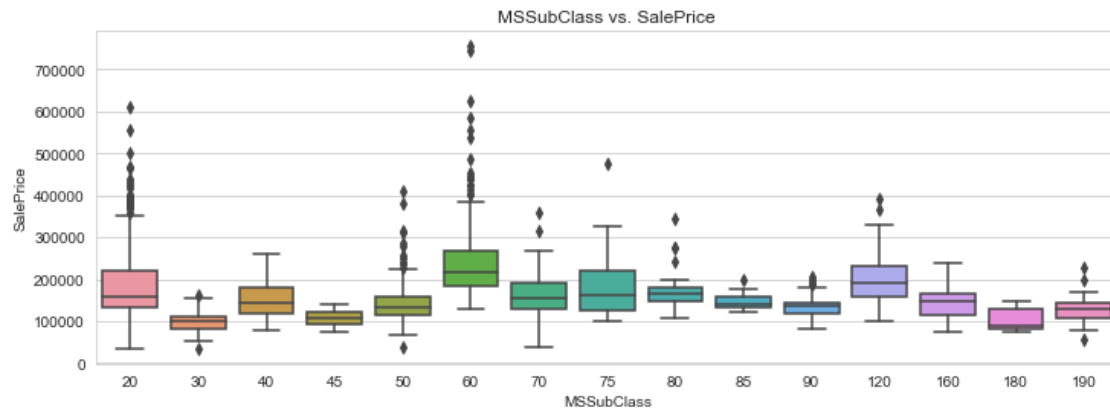


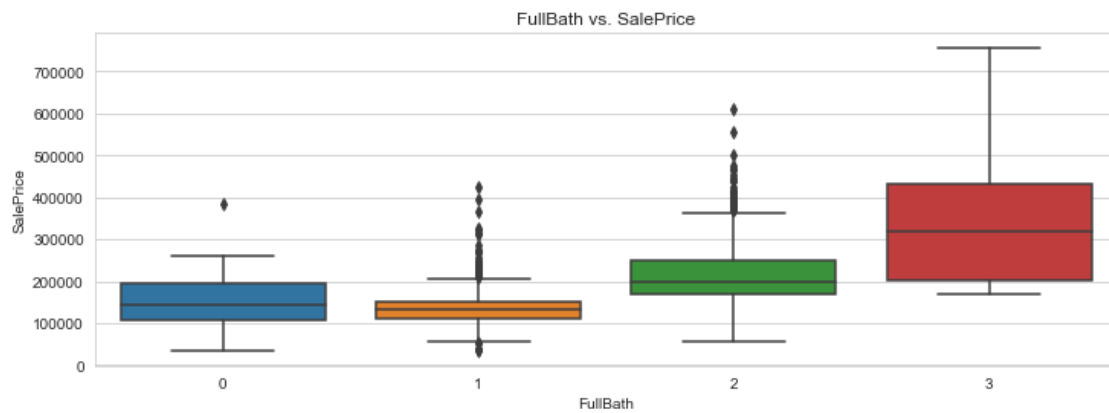
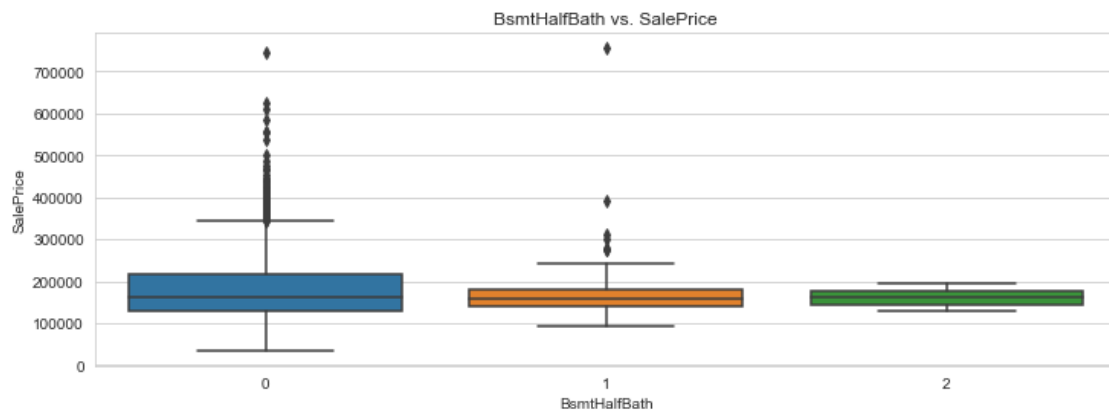
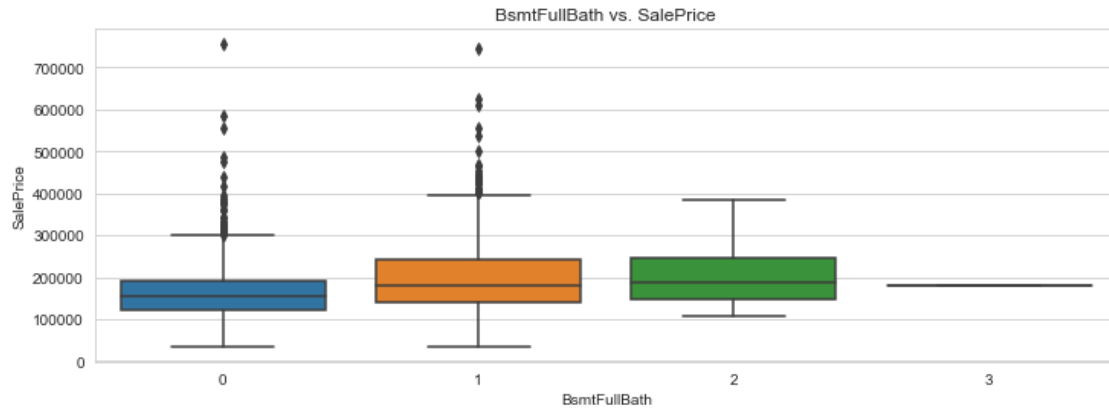


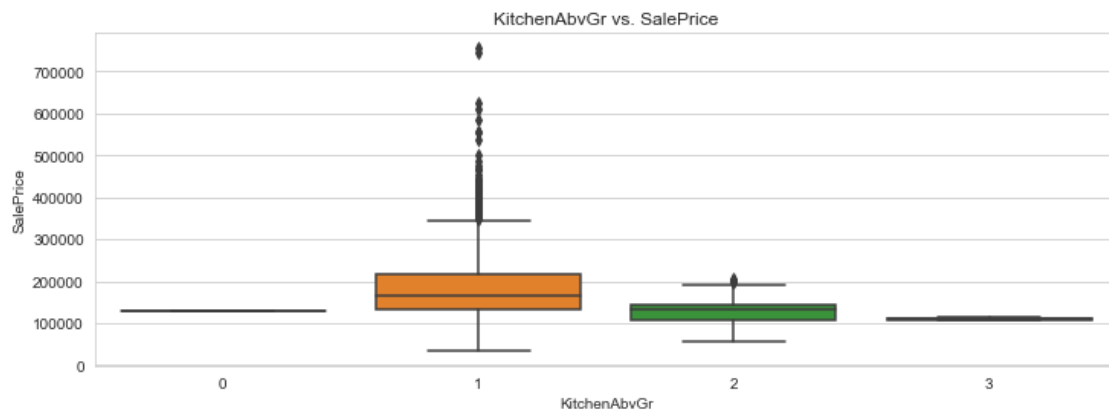
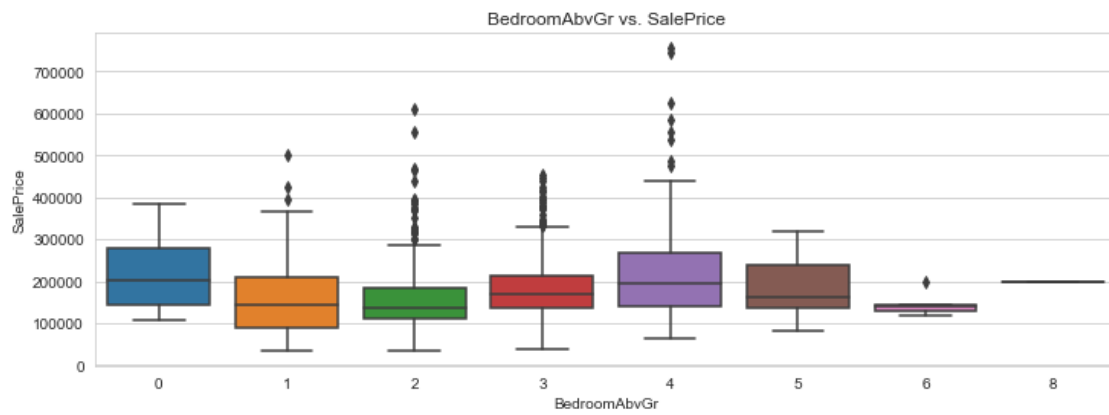
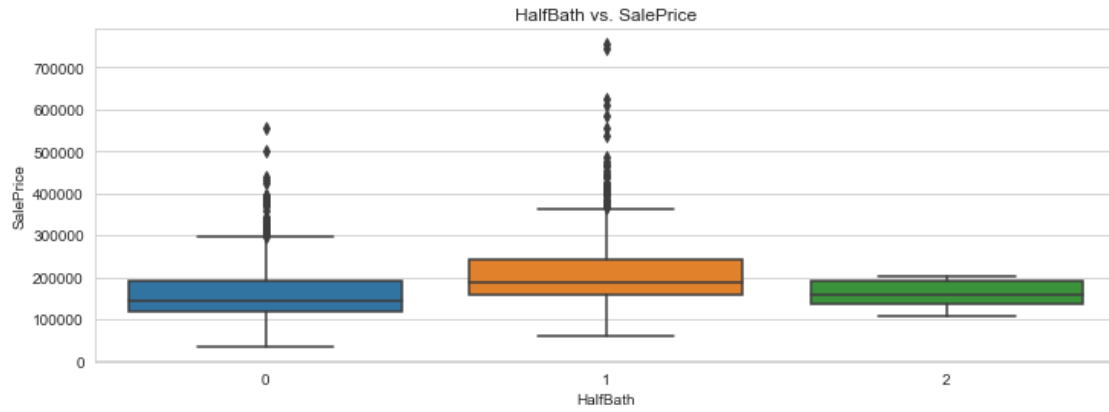


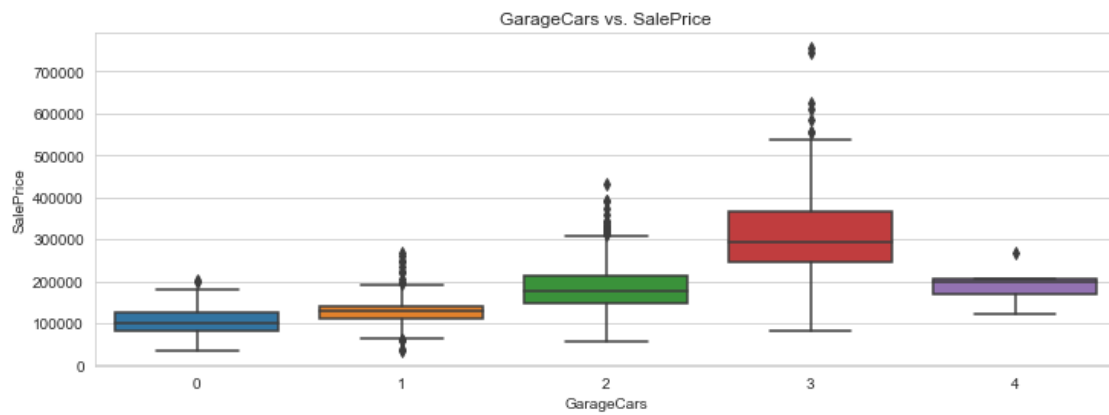
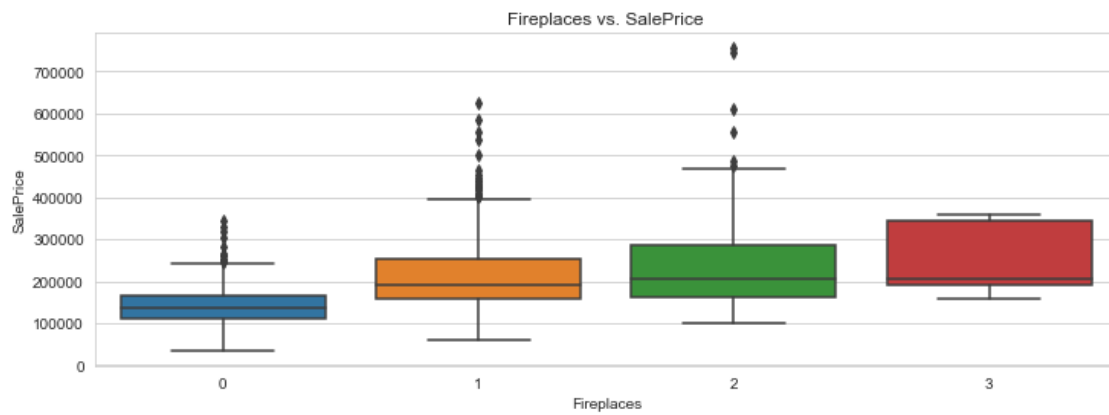
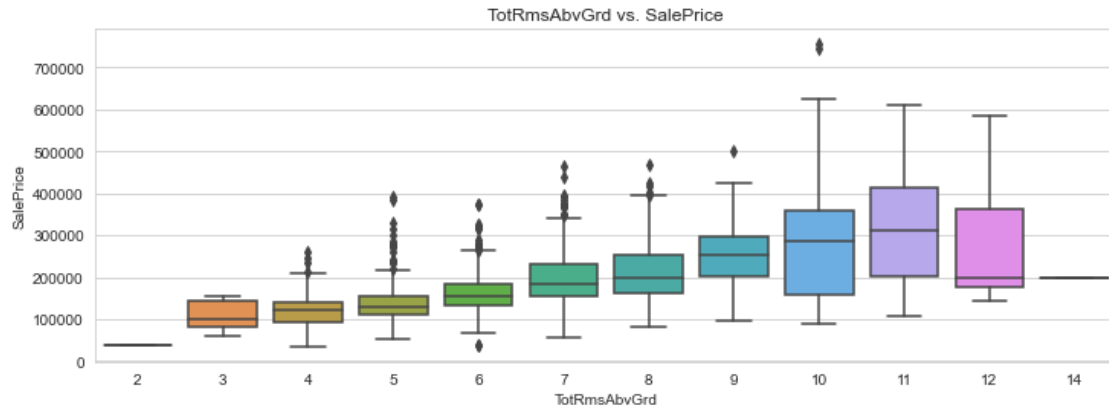
```
[12]: for ax in train[numerical_cat_features]:
      plt.figure(figsize=(12, 4))
      sns.boxplot(x=train[ax], y=train['SalePrice'])
```




```
plt.title(f'{ax} vs. SalePrice')
plt.xlabel(ax)
plt.ylabel('SalePrice')
```















```
[13]: for ax in train[categorical_features]:
plt.figure(figsize=(12, 4))
sns.boxplot(x=train[ax], y=train['SalePrice'], palette='coolwarm')
```




```
plt.title(f'{ax} vs. SalePrice')
plt.xlabel(ax)
plt.ylabel('SalePrice')
```




<ipython-input-13-6794b9677fd5>:2: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).




```
plt.figure(figsize=(12, 4))
```





































































```
[14]: for ax in train[numerical_cont_features]:
    plt.figure(figsize=(12, 4))
    sns.histplot(x=train[ax], bins=30)
    plt.title(f'{ax} distribution')
```




<ipython-input-14-a54a41e995e7>:2: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`).




```
plt.figure(figsize=(12, 4))
```


































```
[15]: for ax in train[numerical_cat_features]:  
      plt.figure(figsize=(12, 4))  
      sns.histplot(x=train[ax], bins=30)
```




```
plt.title(f'{ax} distribution')
```












[16]: # categorical feature distribution represented in boxplots

```
[17]: numerical_cols_w_price = numerical_cols + ['SalePrice']

plt.figure(figsize=(14, 12))
sns.heatmap(train[numerical_cols_w_price].corr())
```


[17]: <AxesSubplot:>



3 Data Cleaning

```
[18]: sns.jointplot(x=train['SalePrice'], y=train['SalePrice'])
```

[18]: <seaborn.axisgrid.JointGrid at 0x1e1895ee9d0>



```
[19]: # method to drop sale price outlier indexes (z score >= 3)

def drop_price_outliers(df):
    i = 0
    drop_indexes = []
    for value in df['SalePrice']:
        if abs((value - df['SalePrice'].mean()) / df['SalePrice'].std()) >= 3:
            drop_indexes.append(i)
        i += 1
    return df.drop(index=drop_indexes).reset_index()
```

```
[20]: train = drop_price_outliers(train)
```

```
[21]: # method to fill numerical cols with sparse data
```

```
def fill_sparse_num_cols(df):  
    for col in df[numerical_cols]:  
        if (df[col].isnull().sum() / len(df[col])) * 100 > 0:  
            df[col] = df[col].fillna(value=round(df[col].mean(), 0))
```

```
[22]: fill_sparse_num_cols(train)  
fill_sparse_num_cols(test)
```

```
[23]: # method to impute values to the features mean with a z score >= 3
```

```
pd.options.mode.chained_assignment = None  
  
def impute_num_outliers(df):  
    z_score_dic = {}  
    i = 1  
    for col in df[numerical_cont_features]:  
        for value in df[col]:  
            if abs((value - df[col].mean()) / df[col].std()) >= 3 and \  
                col not in z_score_dic.keys():  
                z_score_dic[col] = [i - 1]  
            elif abs((value - df[col].mean()) / df[col].std()) >= 3 and \  
                col in z_score_dic.keys():  
                z_score_dic[col].append(i - 1)  
            i += 1  
        if i > len(df[col]):  
            i = 1  
    for key in z_score_dic.keys():  
        for value in z_score_dic[key]:  
            df[key][value] = round(df[key].mean(), 2)  
    return df
```

```
[24]: train = impute_num_outliers(train)  
test = impute_num_outliers(test)
```

```
[25]: # method to drop cols with > 45% missing data
```

```
def drop_sparse_cat_cols(df):  
    drop_cols = []  
    for col in df[categorical_features]:  
        if ((df[col].isnull().sum() / len(df[col])) * 100) > 45:  
            drop_cols.append(col)  
    return df.drop(labels=drop_cols, axis=1)
```

```
[26]: train = drop_sparse_cat_cols(train)  
test = drop_sparse_cat_cols(test)
```

```
[27]: # method to fill missing categorical features with missing object

updated_cat_features = [col for col in train.drop('SalePrice', axis=1) if \
                        train[col].dtype in ['object']]

def fill_missing_cat_cols(df):
    for col in df[updated_cat_features]:
        df[col].fillna(value='Missing', inplace=True)
    return df

[28]: train = fill_missing_cat_cols(train)
test = fill_missing_cat_cols(test)
```

4 Feature Engineering / Scaling

```
[29]: # method to transform categorical data to discrete vars via label encoder

def labeling(df):
    for col in df[updated_cat_features]:
        le = LabelEncoder()
        temp = le.fit_transform(df[col])
        df[f'{col}_labels'] = temp
        df = df.drop(labels=col, axis=1)
    if 'index' in df.columns:
        df = df.drop(labels='index', axis=1)
    return df

[30]: train = labeling(train)
test = labeling(test)

[31]: sns.lineplot(x=train['YrSold'], y=train['SalePrice'])

[31]: <AxesSubplot:xlabel='YrSold', ylabel='SalePrice'>
```




```
[32]: temporal_features = [col for col in train if 'Year' in col or 'Yr' in col]
```

```
[33]: # method to reassign temporal attributes with respect to year sold (dropping_
      ↳ year sold)
```

```
def reassign_temporals(df):
    for col in df[temporal_features]:
        if col != 'YrSold':
            df[col] = df['YrSold'] - df[col]
    return df.drop(labels='YrSold', axis=1)
```

```
[34]: train = reassign_temporals(train)
      test = reassign_temporals(test)
```

```
[35]: # method to normalize data via min max scaler w/ target var
```

```
train_cols = [col for col in train if col != 'Id']
test_cols = [col for col in test if col != 'Id']
```

```
def normalize(df):
    if len(df.columns) == 75:
        cols = train_cols
    else:
        cols = test_cols
    scaler = MinMaxScaler()
```

```
return pd.DataFrame(data=scaler.fit_transform(df[cols]), columns=cols)
```

```
[36]: # train_cols = [col for col in train if col not in ['Id', 'SalePrice']]
# test_cols = [col for col in test if col != 'Id']

# def normalize(df):
#     scaler = MinMaxScaler()
#     if len(df.columns) == 75:
#         df1 = pd.DataFrame(data=scaler.fit_transform(df[train_cols]),
#                               columns=train_cols)
#         df1['SalePrice'] = df['SalePrice']
#         return df1
#     else:
#         return pd.DataFrame(data=scaler.fit_transform(df[test_cols]),
#                               columns=test_cols)
```

```
[37]: train_normal = normalize(train)
test_normal = normalize(test)
```

```
[38]: # method to standardize data via standard scaler w/ target var

def standardize(df):
    if len(df.columns) == 75:
        cols = train_cols
    else:
        cols = test_cols
    scaler = StandardScaler()
    return pd.DataFrame(data=scaler.fit_transform(df[cols]), columns=cols)
```

```
[39]: # def standardize(df):
#     scaler = StandardScaler()
#     if len(df.columns) == 75:
#         df1 = pd.DataFrame(data=scaler.fit_transform(df[train_cols]),
#                               columns=train_cols)
#         df1['SalePrice'] = df['SalePrice']
#         return df1
#     else:
#         return pd.DataFrame(data=scaler.fit_transform(df[test_cols]),
#                               columns=test_cols)
```

```
[40]: train_standard = standardize(train)
test_standard = standardize(test)
```

5 Feature Selection

```
[41]: # feature selection on normalized data

X_train_normal = train_normal.drop('SalePrice', axis=1)
y_train_normal = train_normal['SalePrice']

feature_selection_normal = SelectFromModel(Lasso(alpha=0.005, random_state=0))
feature_selection_normal.fit(X_train_normal, y_train_normal)
```

```
[41]: SelectFromModel(estimator=Lasso(alpha=0.005, random_state=0))
```

```
[42]: selected_feat_normal = X_train_normal.columns[(feature_selection_normal.
↳get_support())]
X_train_normal_selected = X_train_normal[selected_feat_normal]
```

```
[43]: # feature selection on standardized data

X_train_standard = train_standard.drop('SalePrice', axis=1)
y_train_standard = train_standard['SalePrice']

feature_selection_standard = SelectFromModel(Lasso(alpha=0.005, random_state=0))
feature_selection_standard.fit(X_train_standard, y_train_standard)
```

```
[43]: SelectFromModel(estimator=Lasso(alpha=0.005, random_state=0))
```

```
[44]: selected_feat_standard = X_train_standard.columns[(feature_selection_standard.
↳get_support())]
X_train_standard_selected = X_train_standard[selected_feat_standard]
```

```
[45]: count = len(X_train_normal_selected.columns)
for col in X_train_normal_selected.columns:
    if col in X_train_standard_selected.columns:
        count -= 1
print(f'{count} -> same selected cols')
```

0 -> same selected cols

6 Modelling Neural Nets

```
[46]: # train test split overriding normal and standard scaled vars

X_train_normal, X_test_normal, y_train_normal, y_test_normal = \
train_test_split(X_train_normal_selected, y_train_normal, test_size=0.2,
↳random_state=0)
```

```
X_train_standard, X_test_standard, y_train_standard, y_test_standard = \
train_test_split(X_train_standard_selected, y_train_standard, test_size=0.2, \
    ↪random_state=0)
```

[47]: *# configuring neural net w/ adam optimizer for normalized data*

```
# 33 33 1 best

model_normal = Sequential()

model_normal.add(Dense(33, activation='relu'))
model_normal.add(Dropout(0.2))

model_normal.add(Dense(33, activation='relu'))
model_normal.add(Dropout(0.2))

model_normal.add(Dense(1))

model_normal.compile(optimizer='adam', loss='mse')
```

[48]: `model_normal.fit(x=X_train_normal, y=y_train_normal, \`
 `↪validation_data=(X_test_normal, \`
 `↪y_test_normal), \`
 `batch_size=128, epochs=100)`

```
Epoch 1/100
9/9 [=====] - 1s 19ms/step - loss: 0.1353 - val_loss:
0.0523
Epoch 2/100
9/9 [=====] - 0s 4ms/step - loss: 0.0740 - val_loss:
0.0342
Epoch 3/100
9/9 [=====] - 0s 4ms/step - loss: 0.0547 - val_loss:
0.0251
Epoch 4/100
9/9 [=====] - 0s 5ms/step - loss: 0.0456 - val_loss:
0.0213
Epoch 5/100
9/9 [=====] - 0s 5ms/step - loss: 0.0374 - val_loss:
0.0160
Epoch 6/100
9/9 [=====] - 0s 5ms/step - loss: 0.0339 - val_loss:
0.0136
Epoch 7/100
9/9 [=====] - 0s 5ms/step - loss: 0.0292 - val_loss:
0.0123
```

Epoch 8/100
9/9 [=====] - 0s 5ms/step - loss: 0.0283 - val_loss: 0.0100
Epoch 9/100
9/9 [=====] - 0s 5ms/step - loss: 0.0243 - val_loss: 0.0106
Epoch 10/100
9/9 [=====] - 0s 4ms/step - loss: 0.0222 - val_loss: 0.0104
Epoch 11/100
9/9 [=====] - 0s 4ms/step - loss: 0.0228 - val_loss: 0.0107
Epoch 12/100
9/9 [=====] - 0s 4ms/step - loss: 0.0201 - val_loss: 0.0102
Epoch 13/100
9/9 [=====] - 0s 4ms/step - loss: 0.0205 - val_loss: 0.0094
Epoch 14/100
9/9 [=====] - 0s 4ms/step - loss: 0.0189 - val_loss: 0.0097
Epoch 15/100
9/9 [=====] - 0s 5ms/step - loss: 0.0179 - val_loss: 0.0098
Epoch 16/100
9/9 [=====] - 0s 4ms/step - loss: 0.0181 - val_loss: 0.0091
Epoch 17/100
9/9 [=====] - 0s 4ms/step - loss: 0.0186 - val_loss: 0.0093
Epoch 18/100
9/9 [=====] - 0s 4ms/step - loss: 0.0168 - val_loss: 0.0095
Epoch 19/100
9/9 [=====] - 0s 5ms/step - loss: 0.0160 - val_loss: 0.0083
Epoch 20/100
9/9 [=====] - 0s 4ms/step - loss: 0.0159 - val_loss: 0.0100
Epoch 21/100
9/9 [=====] - 0s 4ms/step - loss: 0.0151 - val_loss: 0.0095
Epoch 22/100
9/9 [=====] - 0s 4ms/step - loss: 0.0159 - val_loss: 0.0079
Epoch 23/100
9/9 [=====] - 0s 4ms/step - loss: 0.0156 - val_loss: 0.0092

Epoch 24/100
9/9 [=====] - 0s 4ms/step - loss: 0.0149 - val_loss: 0.0093
Epoch 25/100
9/9 [=====] - 0s 4ms/step - loss: 0.0135 - val_loss: 0.0081
Epoch 26/100
9/9 [=====] - 0s 4ms/step - loss: 0.0138 - val_loss: 0.0086
Epoch 27/100
9/9 [=====] - 0s 5ms/step - loss: 0.0139 - val_loss: 0.0082
Epoch 28/100
9/9 [=====] - 0s 4ms/step - loss: 0.0141 - val_loss: 0.0085
Epoch 29/100
9/9 [=====] - 0s 4ms/step - loss: 0.0125 - val_loss: 0.0079
Epoch 30/100
9/9 [=====] - 0s 4ms/step - loss: 0.0129 - val_loss: 0.0093
Epoch 31/100
9/9 [=====] - 0s 4ms/step - loss: 0.0127 - val_loss: 0.0081
Epoch 32/100
9/9 [=====] - 0s 4ms/step - loss: 0.0122 - val_loss: 0.0085
Epoch 33/100
9/9 [=====] - 0s 4ms/step - loss: 0.0130 - val_loss: 0.0073
Epoch 34/100
9/9 [=====] - 0s 4ms/step - loss: 0.0115 - val_loss: 0.0092
Epoch 35/100
9/9 [=====] - 0s 4ms/step - loss: 0.0128 - val_loss: 0.0078
Epoch 36/100
9/9 [=====] - 0s 4ms/step - loss: 0.0121 - val_loss: 0.0077
Epoch 37/100
9/9 [=====] - 0s 4ms/step - loss: 0.0115 - val_loss: 0.0080
Epoch 38/100
9/9 [=====] - 0s 4ms/step - loss: 0.0117 - val_loss: 0.0083
Epoch 39/100
9/9 [=====] - 0s 4ms/step - loss: 0.0110 - val_loss: 0.0070

Epoch 40/100
9/9 [=====] - 0s 4ms/step - loss: 0.0109 - val_loss: 0.0074
Epoch 41/100
9/9 [=====] - 0s 4ms/step - loss: 0.0110 - val_loss: 0.0079
Epoch 42/100
9/9 [=====] - 0s 4ms/step - loss: 0.0115 - val_loss: 0.0076
Epoch 43/100
9/9 [=====] - 0s 4ms/step - loss: 0.0114 - val_loss: 0.0076
Epoch 44/100
9/9 [=====] - 0s 4ms/step - loss: 0.0104 - val_loss: 0.0074
Epoch 45/100
9/9 [=====] - 0s 4ms/step - loss: 0.0100 - val_loss: 0.0076
Epoch 46/100
9/9 [=====] - 0s 4ms/step - loss: 0.0103 - val_loss: 0.0077
Epoch 47/100
9/9 [=====] - 0s 4ms/step - loss: 0.0107 - val_loss: 0.0066
Epoch 48/100
9/9 [=====] - 0s 4ms/step - loss: 0.0103 - val_loss: 0.0073
Epoch 49/100
9/9 [=====] - 0s 4ms/step - loss: 0.0101 - val_loss: 0.0075
Epoch 50/100
9/9 [=====] - 0s 4ms/step - loss: 0.0100 - val_loss: 0.0066
Epoch 51/100
9/9 [=====] - 0s 4ms/step - loss: 0.0105 - val_loss: 0.0073
Epoch 52/100
9/9 [=====] - 0s 4ms/step - loss: 0.0097 - val_loss: 0.0067
Epoch 53/100
9/9 [=====] - 0s 4ms/step - loss: 0.0089 - val_loss: 0.0071
Epoch 54/100
9/9 [=====] - 0s 4ms/step - loss: 0.0100 - val_loss: 0.0072
Epoch 55/100
9/9 [=====] - 0s 4ms/step - loss: 0.0099 - val_loss: 0.0065

Epoch 56/100
9/9 [=====] - 0s 4ms/step - loss: 0.0099 - val_loss: 0.0073
Epoch 57/100
9/9 [=====] - 0s 4ms/step - loss: 0.0085 - val_loss: 0.0064
Epoch 58/100
9/9 [=====] - 0s 4ms/step - loss: 0.0091 - val_loss: 0.0068
Epoch 59/100
9/9 [=====] - 0s 4ms/step - loss: 0.0087 - val_loss: 0.0059
Epoch 60/100
9/9 [=====] - 0s 4ms/step - loss: 0.0087 - val_loss: 0.0074
Epoch 61/100
9/9 [=====] - 0s 4ms/step - loss: 0.0093 - val_loss: 0.0062
Epoch 62/100
9/9 [=====] - 0s 5ms/step - loss: 0.0095 - val_loss: 0.0069
Epoch 63/100
9/9 [=====] - 0s 5ms/step - loss: 0.0091 - val_loss: 0.0069
Epoch 64/100
9/9 [=====] - 0s 5ms/step - loss: 0.0092 - val_loss: 0.0065
Epoch 65/100
9/9 [=====] - 0s 5ms/step - loss: 0.0088 - val_loss: 0.0070
Epoch 66/100
9/9 [=====] - 0s 4ms/step - loss: 0.0083 - val_loss: 0.0066
Epoch 67/100
9/9 [=====] - 0s 4ms/step - loss: 0.0089 - val_loss: 0.0063
Epoch 68/100
9/9 [=====] - 0s 4ms/step - loss: 0.0085 - val_loss: 0.0064
Epoch 69/100
9/9 [=====] - 0s 4ms/step - loss: 0.0085 - val_loss: 0.0066
Epoch 70/100
9/9 [=====] - 0s 4ms/step - loss: 0.0085 - val_loss: 0.0064
Epoch 71/100
9/9 [=====] - 0s 4ms/step - loss: 0.0083 - val_loss: 0.0067

Epoch 72/100
9/9 [=====] - 0s 4ms/step - loss: 0.0084 - val_loss: 0.0060
Epoch 73/100
9/9 [=====] - 0s 4ms/step - loss: 0.0082 - val_loss: 0.0067
Epoch 74/100
9/9 [=====] - 0s 4ms/step - loss: 0.0089 - val_loss: 0.0072
Epoch 75/100
9/9 [=====] - 0s 4ms/step - loss: 0.0084 - val_loss: 0.0060
Epoch 76/100
9/9 [=====] - 0s 4ms/step - loss: 0.0086 - val_loss: 0.0069
Epoch 77/100
9/9 [=====] - 0s 4ms/step - loss: 0.0076 - val_loss: 0.0056
Epoch 78/100
9/9 [=====] - 0s 4ms/step - loss: 0.0077 - val_loss: 0.0072
Epoch 79/100
9/9 [=====] - 0s 4ms/step - loss: 0.0082 - val_loss: 0.0062
Epoch 80/100
9/9 [=====] - 0s 4ms/step - loss: 0.0079 - val_loss: 0.0059
Epoch 81/100
9/9 [=====] - 0s 4ms/step - loss: 0.0083 - val_loss: 0.0067
Epoch 82/100
9/9 [=====] - 0s 4ms/step - loss: 0.0083 - val_loss: 0.0058
Epoch 83/100
9/9 [=====] - 0s 4ms/step - loss: 0.0079 - val_loss: 0.0061
Epoch 84/100
9/9 [=====] - 0s 4ms/step - loss: 0.0080 - val_loss: 0.0062
Epoch 85/100
9/9 [=====] - 0s 4ms/step - loss: 0.0085 - val_loss: 0.0063
Epoch 86/100
9/9 [=====] - 0s 4ms/step - loss: 0.0080 - val_loss: 0.0067
Epoch 87/100
9/9 [=====] - 0s 4ms/step - loss: 0.0073 - val_loss: 0.0057

```


Epoch 88/100
9/9 [=====] - 0s 4ms/step - loss: 0.0078 - val_loss: 0.0067
Epoch 89/100
9/9 [=====] - 0s 4ms/step - loss: 0.0080 - val_loss: 0.0064
Epoch 90/100
9/9 [=====] - 0s 4ms/step - loss: 0.0069 - val_loss: 0.0059
Epoch 91/100
9/9 [=====] - 0s 4ms/step - loss: 0.0078 - val_loss: 0.0070
Epoch 92/100
9/9 [=====] - 0s 4ms/step - loss: 0.0079 - val_loss: 0.0058
Epoch 93/100
9/9 [=====] - 0s 4ms/step - loss: 0.0076 - val_loss: 0.0061
Epoch 94/100
9/9 [=====] - 0s 4ms/step - loss: 0.0073 - val_loss: 0.0061
Epoch 95/100
9/9 [=====] - 0s 4ms/step - loss: 0.0079 - val_loss: 0.0071
Epoch 96/100
9/9 [=====] - 0s 5ms/step - loss: 0.0076 - val_loss: 0.0056
Epoch 97/100
9/9 [=====] - 0s 4ms/step - loss: 0.0071 - val_loss: 0.0062
Epoch 98/100
9/9 [=====] - 0s 4ms/step - loss: 0.0076 - val_loss: 0.0063
Epoch 99/100
9/9 [=====] - 0s 4ms/step - loss: 0.0074 - val_loss: 0.0059
Epoch 100/100
9/9 [=====] - 0s 4ms/step - loss: 0.0074 - val_loss: 0.0062

```

```
[48]: <keras.callbacks.History at 0x1e18b23c6a0>
```

```
[49]: normal_losses = pd.DataFrame(model_normal.history.history)
normal_losses.plot()
```

```
[49]: <AxesSubplot:>
```



```
[50]: normal_predictions = model_normal.predict(X_test_normal)
```

```
[51]: mean_absolute_error(y_test_normal, normal_predictions)
```

```
[51]: 0.054884360017891126
```

```
[52]: np.sqrt(mean_squared_error(y_test_normal, normal_predictions))
```


```
[52]: 0.07872164979893682
```

```
[53]: explained_variance_score(y_test_normal, normal_predictions)
```

```
[53]: 0.8412406475412091
```

```
[54]: plt.scatter(y_test_normal, normal_predictions)
plt.plot(y_test_normal, y_test_normal, 'r')
```

```
[54]: [<matplotlib.lines.Line2D at 0x1e18eb476a0>]
```



```
[55]: # configuring neural net w/ adam optimizer for standardized data
```

```
# 33 33 1 best
```

```
model_standard = Sequential()

model_standard.add(Dense(33, activation='relu'))
model_standard.add(Dropout(0.2))

model_standard.add(Dense(33, activation='relu'))
model_standard.add(Dropout(0.2))

model_standard.add(Dense(1))

model_standard.compile(optimizer='adam', loss='mse')
```

```
[56]: model_standard.fit(x=X_train_standard, y=y_train_standard, \
    ↪ validation_data=(X_test_standard, \
    ↪ y_test_standard), \
    batch_size=128, epochs=100)
```

Epoch 1/100

9/9 [=====] - 0s 16ms/step - loss: 1.2218 - val_loss: 0.6603

Epoch 2/100

9/9 [=====] - 0s 4ms/step - loss: 0.8315 - val_loss: 0.3434
Epoch 3/100
9/9 [=====] - 0s 4ms/step - loss: 0.5968 - val_loss: 0.2691
Epoch 4/100
9/9 [=====] - 0s 4ms/step - loss: 0.5301 - val_loss: 0.2488
Epoch 5/100
9/9 [=====] - 0s 4ms/step - loss: 0.4791 - val_loss: 0.2309
Epoch 6/100
9/9 [=====] - 0s 4ms/step - loss: 0.4046 - val_loss: 0.2138
Epoch 7/100
9/9 [=====] - 0s 4ms/step - loss: 0.3954 - val_loss: 0.2011
Epoch 8/100
9/9 [=====] - 0s 4ms/step - loss: 0.4167 - val_loss: 0.2000
Epoch 9/100
9/9 [=====] - 0s 4ms/step - loss: 0.3740 - val_loss: 0.1949
Epoch 10/100
9/9 [=====] - 0s 4ms/step - loss: 0.3390 - val_loss: 0.1867
Epoch 11/100
9/9 [=====] - 0s 4ms/step - loss: 0.3009 - val_loss: 0.1866
Epoch 12/100
9/9 [=====] - 0s 4ms/step - loss: 0.3248 - val_loss: 0.1791
Epoch 13/100
9/9 [=====] - 0s 4ms/step - loss: 0.2865 - val_loss: 0.1649
Epoch 14/100
9/9 [=====] - 0s 4ms/step - loss: 0.2967 - val_loss: 0.1591
Epoch 15/100
9/9 [=====] - 0s 4ms/step - loss: 0.3163 - val_loss: 0.1573
Epoch 16/100
9/9 [=====] - 0s 4ms/step - loss: 0.2708 - val_loss: 0.1584
Epoch 17/100
9/9 [=====] - 0s 5ms/step - loss: 0.3058 - val_loss: 0.1572
Epoch 18/100

9/9 [=====] - 0s 5ms/step - loss: 0.2642 - val_loss:
0.1554
Epoch 19/100
9/9 [=====] - 0s 4ms/step - loss: 0.2757 - val_loss:
0.1544
Epoch 20/100
9/9 [=====] - 0s 4ms/step - loss: 0.2556 - val_loss:
0.1546
Epoch 21/100
9/9 [=====] - 0s 4ms/step - loss: 0.2115 - val_loss:
0.1484
Epoch 22/100
9/9 [=====] - 0s 4ms/step - loss: 0.2239 - val_loss:
0.1425
Epoch 23/100
9/9 [=====] - 0s 4ms/step - loss: 0.2365 - val_loss:
0.1406
Epoch 24/100
9/9 [=====] - 0s 4ms/step - loss: 0.2252 - val_loss:
0.1408
Epoch 25/100
9/9 [=====] - 0s 4ms/step - loss: 0.2205 - val_loss:
0.1408
Epoch 26/100
9/9 [=====] - 0s 4ms/step - loss: 0.2258 - val_loss:
0.1491
Epoch 27/100
9/9 [=====] - 0s 4ms/step - loss: 0.2308 - val_loss:
0.1570
Epoch 28/100
9/9 [=====] - 0s 4ms/step - loss: 0.2156 - val_loss:
0.1456
Epoch 29/100
9/9 [=====] - 0s 4ms/step - loss: 0.2009 - val_loss:
0.1379
Epoch 30/100
9/9 [=====] - 0s 4ms/step - loss: 0.1909 - val_loss:
0.1440
Epoch 31/100
9/9 [=====] - 0s 4ms/step - loss: 0.1900 - val_loss:
0.1501
Epoch 32/100
9/9 [=====] - 0s 4ms/step - loss: 0.2193 - val_loss:
0.1405
Epoch 33/100
9/9 [=====] - 0s 4ms/step - loss: 0.1898 - val_loss:
0.1304
Epoch 34/100

9/9 [=====] - 0s 4ms/step - loss: 0.1910 - val_loss:
0.1362
Epoch 35/100
9/9 [=====] - 0s 4ms/step - loss: 0.1826 - val_loss:
0.1418
Epoch 36/100
9/9 [=====] - 0s 4ms/step - loss: 0.1910 - val_loss:
0.1354
Epoch 37/100
9/9 [=====] - 0s 4ms/step - loss: 0.1794 - val_loss:
0.1436
Epoch 38/100
9/9 [=====] - 0s 4ms/step - loss: 0.1756 - val_loss:
0.1460
Epoch 39/100
9/9 [=====] - 0s 4ms/step - loss: 0.1787 - val_loss:
0.1375
Epoch 40/100
9/9 [=====] - 0s 4ms/step - loss: 0.1653 - val_loss:
0.1299
Epoch 41/100
9/9 [=====] - 0s 4ms/step - loss: 0.1886 - val_loss:
0.1339
Epoch 42/100
9/9 [=====] - 0s 4ms/step - loss: 0.1622 - val_loss:
0.1416
Epoch 43/100
9/9 [=====] - 0s 4ms/step - loss: 0.1595 - val_loss:
0.1317
Epoch 44/100
9/9 [=====] - 0s 4ms/step - loss: 0.1623 - val_loss:
0.1288
Epoch 45/100
9/9 [=====] - 0s 4ms/step - loss: 0.1570 - val_loss:
0.1328
Epoch 46/100
9/9 [=====] - 0s 4ms/step - loss: 0.1595 - val_loss:
0.1375
Epoch 47/100
9/9 [=====] - 0s 4ms/step - loss: 0.1692 - val_loss:
0.1398
Epoch 48/100
9/9 [=====] - 0s 4ms/step - loss: 0.1633 - val_loss:
0.1362
Epoch 49/100
9/9 [=====] - 0s 4ms/step - loss: 0.1509 - val_loss:
0.1299
Epoch 50/100

9/9 [=====] - 0s 4ms/step - loss: 0.1530 - val_loss:
0.1320
Epoch 51/100
9/9 [=====] - 0s 4ms/step - loss: 0.1390 - val_loss:
0.1331
Epoch 52/100
9/9 [=====] - 0s 4ms/step - loss: 0.1502 - val_loss:
0.1330
Epoch 53/100
9/9 [=====] - 0s 4ms/step - loss: 0.1483 - val_loss:
0.1350
Epoch 54/100
9/9 [=====] - 0s 4ms/step - loss: 0.1345 - val_loss:
0.1361
Epoch 55/100
9/9 [=====] - 0s 4ms/step - loss: 0.1428 - val_loss:
0.1367
Epoch 56/100
9/9 [=====] - 0s 4ms/step - loss: 0.1540 - val_loss:
0.1320
Epoch 57/100
9/9 [=====] - 0s 4ms/step - loss: 0.1430 - val_loss:
0.1286
Epoch 58/100
9/9 [=====] - 0s 4ms/step - loss: 0.1442 - val_loss:
0.1404
Epoch 59/100
9/9 [=====] - 0s 4ms/step - loss: 0.1416 - val_loss:
0.1439
Epoch 60/100
9/9 [=====] - 0s 4ms/step - loss: 0.1493 - val_loss:
0.1284
Epoch 61/100
9/9 [=====] - 0s 4ms/step - loss: 0.1461 - val_loss:
0.1224
Epoch 62/100
9/9 [=====] - 0s 4ms/step - loss: 0.1331 - val_loss:
0.1351
Epoch 63/100
9/9 [=====] - 0s 4ms/step - loss: 0.1399 - val_loss:
0.1425
Epoch 64/100
9/9 [=====] - 0s 4ms/step - loss: 0.1423 - val_loss:
0.1305
Epoch 65/100
9/9 [=====] - 0s 4ms/step - loss: 0.1332 - val_loss:
0.1262
Epoch 66/100


```
9/9 [=====] - 0s 4ms/step - loss: 0.1450 - val_loss: 0.1369
Epoch 67/100
9/9 [=====] - 0s 4ms/step - loss: 0.1271 - val_loss: 0.1402
Epoch 68/100
9/9 [=====] - 0s 4ms/step - loss: 0.1313 - val_loss: 0.1337
Epoch 69/100
9/9 [=====] - 0s 4ms/step - loss: 0.1265 - val_loss: 0.1255
Epoch 70/100
9/9 [=====] - 0s 4ms/step - loss: 0.1275 - val_loss: 0.1284
Epoch 71/100
9/9 [=====] - 0s 4ms/step - loss: 0.1330 - val_loss: 0.1421
Epoch 72/100
9/9 [=====] - 0s 4ms/step - loss: 0.1292 - val_loss: 0.1374
Epoch 73/100
9/9 [=====] - 0s 4ms/step - loss: 0.1296 - val_loss: 0.1280
Epoch 74/100
9/9 [=====] - 0s 4ms/step - loss: 0.1316 - val_loss: 0.1319
Epoch 75/100
9/9 [=====] - 0s 4ms/step - loss: 0.1232 - val_loss: 0.1367
Epoch 76/100
9/9 [=====] - 0s 4ms/step - loss: 0.1273 - val_loss: 0.1302
Epoch 77/100
9/9 [=====] - 0s 4ms/step - loss: 0.1204 - val_loss: 0.1366
Epoch 78/100
9/9 [=====] - 0s 4ms/step - loss: 0.1269 - val_loss: 0.1355
Epoch 79/100
9/9 [=====] - 0s 4ms/step - loss: 0.1163 - val_loss: 0.1296
Epoch 80/100
9/9 [=====] - 0s 4ms/step - loss: 0.1170 - val_loss: 0.1257
Epoch 81/100
9/9 [=====] - 0s 4ms/step - loss: 0.1331 - val_loss: 0.1346
Epoch 82/100
```

9/9 [=====] - 0s 4ms/step - loss: 0.1249 - val_loss: 0.1334
Epoch 83/100
9/9 [=====] - 0s 4ms/step - loss: 0.1223 - val_loss: 0.1264
Epoch 84/100
9/9 [=====] - 0s 5ms/step - loss: 0.1169 - val_loss: 0.1267
Epoch 85/100
9/9 [=====] - 0s 5ms/step - loss: 0.1132 - val_loss: 0.1327
Epoch 86/100
9/9 [=====] - 0s 5ms/step - loss: 0.1361 - val_loss: 0.1282
Epoch 87/100
9/9 [=====] - 0s 4ms/step - loss: 0.1159 - val_loss: 0.1336
Epoch 88/100
9/9 [=====] - 0s 4ms/step - loss: 0.1217 - val_loss: 0.1375
Epoch 89/100
9/9 [=====] - 0s 4ms/step - loss: 0.1019 - val_loss: 0.1327
Epoch 90/100
9/9 [=====] - 0s 4ms/step - loss: 0.1097 - val_loss: 0.1326
Epoch 91/100
9/9 [=====] - 0s 4ms/step - loss: 0.1107 - val_loss: 0.1295
Epoch 92/100
9/9 [=====] - 0s 4ms/step - loss: 0.1238 - val_loss: 0.1328
Epoch 93/100
9/9 [=====] - 0s 4ms/step - loss: 0.1193 - val_loss: 0.1349
Epoch 94/100
9/9 [=====] - 0s 4ms/step - loss: 0.1139 - val_loss: 0.1308
Epoch 95/100
9/9 [=====] - 0s 4ms/step - loss: 0.1102 - val_loss: 0.1247
Epoch 96/100
9/9 [=====] - 0s 4ms/step - loss: 0.1070 - val_loss: 0.1335
Epoch 97/100
9/9 [=====] - 0s 4ms/step - loss: 0.1176 - val_loss: 0.1383
Epoch 98/100

```


9/9 [=====] - 0s 4ms/step - loss: 0.1037 - val_loss:
0.1282
Epoch 99/100
9/9 [=====] - 0s 4ms/step - loss: 0.1083 - val_loss:
0.1217
Epoch 100/100
9/9 [=====] - 0s 4ms/step - loss: 0.1229 - val_loss:
0.1278

```

```
[56]: <keras.callbacks.History at 0x1e18ac423a0>
```

```
[57]: standard_losses = pd.DataFrame(model_standard.history.history)
standard_losses.plot()
```

```
[57]: <AxesSubplot:>
```



```
[58]: standard_predictions = model_standard.predict(X_test_standard)
```

```
[59]: mean_absolute_error(y_test_standard, standard_predictions)
```

```
[59]: 0.24173186557614557
```

```
[60]: np.sqrt(mean_squared_error(y_test_standard, standard_predictions))
```


```
[60]: 0.3575448254640191
```

```
[61]: explained_variance_score(y_test_standard, standard_predictions)
```

```
[61]: 0.8870123453564851
```

```
[62]: plt.scatter(y_test_standard, standard_predictions)
plt.plot(y_test_standard, y_test_standard, 'r')
```

```
[62]: [<matplotlib.lines.Line2D at 0x1e18b1b5e80>]
```



```
[63]: # TODO: neural net on full data set (maybe)
# TODO: neural net on pca of standardized data (non-lasso-selected?) (maybe)
```

```
[64]: # modelling neural net for all cols
```

```
X_train_standard = train_standard.drop('SalePrice', axis=1)
y_train_standard = train_standard['SalePrice']
```

```
X_train_stand, X_test_stand, y_train_stand, y_test_stand = \
    train_test_split(X_train_standard, \
```

```
        y_train_standard, \
```

```
        test_size=0.2, \
```

```
        random_state=0)
```

[65]: *# configuring neural net w/ adam optimizer for standardized non-selected data*

```
# rmse:
# 147 77 39 .0642
# 147 147 77 39 1 .0648
# 147 77 77 1 .0706
# 147 147 77 .0731
# 147 77 77 .0741
# 147 147 1 .0777

model_stand = Sequential()

model_stand.add(Dense(147, activation='relu'))
model_stand.add(Dropout(0.2))

# model_stand.add(Dense(147, activation='relu'))
# model_stand.add(Dropout(0.2))

model_stand.add(Dense(77, activation='relu'))
model_stand.add(Dropout(0.2))

model_stand.add(Dense(39, activation='relu'))
model_stand.add(Dropout(0.2))

model_stand.add(Dense(1))

model_stand.compile(optimizer='adam', loss='mse')
```

[66]: `model_stand.fit(x=X_train_stand, y=y_train_stand, \`
 `↪ validation_data=(X_test_stand, \`
 `↪ y_test_stand), \`
 `batch_size=128, epochs=100)`

```
Epoch 1/100
9/9 [=====] - 1s 18ms/step - loss: 0.8462 - val_loss:
0.3582
Epoch 2/100
9/9 [=====] - 0s 6ms/step - loss: 0.4516 - val_loss:
0.2216
Epoch 3/100
9/9 [=====] - 0s 5ms/step - loss: 0.3519 - val_loss:
0.2334
Epoch 4/100
9/9 [=====] - 0s 5ms/step - loss: 0.3193 - val_loss:
0.2155
Epoch 5/100
```

9/9 [=====] - 0s 6ms/step - loss: 0.2602 - val_loss:
0.1810
Epoch 6/100
9/9 [=====] - 0s 5ms/step - loss: 0.2498 - val_loss:
0.1775
Epoch 7/100
9/9 [=====] - 0s 6ms/step - loss: 0.2220 - val_loss:
0.1788
Epoch 8/100
9/9 [=====] - 0s 5ms/step - loss: 0.2100 - val_loss:
0.1741
Epoch 9/100
9/9 [=====] - 0s 5ms/step - loss: 0.2102 - val_loss:
0.1702
Epoch 10/100
9/9 [=====] - 0s 5ms/step - loss: 0.2067 - val_loss:
0.1711
Epoch 11/100
9/9 [=====] - 0s 6ms/step - loss: 0.2148 - val_loss:
0.1664
Epoch 12/100
9/9 [=====] - 0s 5ms/step - loss: 0.1970 - val_loss:
0.1555
Epoch 13/100
9/9 [=====] - 0s 6ms/step - loss: 0.1972 - val_loss:
0.1476
Epoch 14/100
9/9 [=====] - 0s 5ms/step - loss: 0.1661 - val_loss:
0.1438
Epoch 15/100
9/9 [=====] - 0s 5ms/step - loss: 0.1661 - val_loss:
0.1487
Epoch 16/100
9/9 [=====] - 0s 6ms/step - loss: 0.1622 - val_loss:
0.1444
Epoch 17/100
9/9 [=====] - 0s 6ms/step - loss: 0.1573 - val_loss:
0.1407
Epoch 18/100
9/9 [=====] - 0s 6ms/step - loss: 0.1531 - val_loss:
0.1468
Epoch 19/100
9/9 [=====] - 0s 6ms/step - loss: 0.1393 - val_loss:
0.1474
Epoch 20/100
9/9 [=====] - 0s 6ms/step - loss: 0.1440 - val_loss:
0.1513
Epoch 21/100

9/9 [=====] - 0s 5ms/step - loss: 0.1326 - val_loss:
0.1279
Epoch 22/100
9/9 [=====] - 0s 6ms/step - loss: 0.1433 - val_loss:
0.1633
Epoch 23/100
9/9 [=====] - 0s 5ms/step - loss: 0.1425 - val_loss:
0.1403
Epoch 24/100
9/9 [=====] - 0s 5ms/step - loss: 0.1268 - val_loss:
0.1389
Epoch 25/100
9/9 [=====] - 0s 5ms/step - loss: 0.1367 - val_loss:
0.1464
Epoch 26/100
9/9 [=====] - 0s 6ms/step - loss: 0.1164 - val_loss:
0.1342
Epoch 27/100
9/9 [=====] - 0s 5ms/step - loss: 0.1280 - val_loss:
0.1475
Epoch 28/100
9/9 [=====] - 0s 6ms/step - loss: 0.1277 - val_loss:
0.1337
Epoch 29/100
9/9 [=====] - 0s 5ms/step - loss: 0.1170 - val_loss:
0.1343
Epoch 30/100
9/9 [=====] - 0s 6ms/step - loss: 0.1100 - val_loss:
0.1374
Epoch 31/100
9/9 [=====] - 0s 5ms/step - loss: 0.1021 - val_loss:
0.1284
Epoch 32/100
9/9 [=====] - 0s 5ms/step - loss: 0.1289 - val_loss:
0.1519
Epoch 33/100
9/9 [=====] - 0s 5ms/step - loss: 0.1123 - val_loss:
0.1290
Epoch 34/100
9/9 [=====] - 0s 5ms/step - loss: 0.1130 - val_loss:
0.1419
Epoch 35/100
9/9 [=====] - 0s 6ms/step - loss: 0.1042 - val_loss:
0.1542
Epoch 36/100
9/9 [=====] - 0s 5ms/step - loss: 0.0990 - val_loss:
0.1332
Epoch 37/100

9/9 [=====] - 0s 5ms/step - loss: 0.1056 - val_loss: 0.1356
Epoch 38/100
9/9 [=====] - 0s 5ms/step - loss: 0.1021 - val_loss: 0.1333
Epoch 39/100
9/9 [=====] - 0s 6ms/step - loss: 0.1019 - val_loss: 0.1346
Epoch 40/100
9/9 [=====] - 0s 5ms/step - loss: 0.1040 - val_loss: 0.1288
Epoch 41/100
9/9 [=====] - 0s 5ms/step - loss: 0.0876 - val_loss: 0.1365
Epoch 42/100
9/9 [=====] - 0s 6ms/step - loss: 0.0945 - val_loss: 0.1454
Epoch 43/100
9/9 [=====] - 0s 6ms/step - loss: 0.0994 - val_loss: 0.1292
Epoch 44/100
9/9 [=====] - 0s 5ms/step - loss: 0.0968 - val_loss: 0.1346
Epoch 45/100
9/9 [=====] - 0s 5ms/step - loss: 0.0943 - val_loss: 0.1313
Epoch 46/100
9/9 [=====] - 0s 5ms/step - loss: 0.0967 - val_loss: 0.1423
Epoch 47/100
9/9 [=====] - 0s 5ms/step - loss: 0.0876 - val_loss: 0.1397
Epoch 48/100
9/9 [=====] - 0s 5ms/step - loss: 0.0932 - val_loss: 0.1299
Epoch 49/100
9/9 [=====] - 0s 5ms/step - loss: 0.0851 - val_loss: 0.1317
Epoch 50/100
9/9 [=====] - 0s 5ms/step - loss: 0.0805 - val_loss: 0.1286
Epoch 51/100
9/9 [=====] - 0s 5ms/step - loss: 0.0829 - val_loss: 0.1346
Epoch 52/100
9/9 [=====] - 0s 5ms/step - loss: 0.0847 - val_loss: 0.1411
Epoch 53/100

9/9 [=====] - 0s 5ms/step - loss: 0.0986 - val_loss:
0.1297
Epoch 54/100
9/9 [=====] - 0s 5ms/step - loss: 0.0898 - val_loss:
0.1396
Epoch 55/100
9/9 [=====] - 0s 5ms/step - loss: 0.0908 - val_loss:
0.1393
Epoch 56/100
9/9 [=====] - 0s 6ms/step - loss: 0.0788 - val_loss:
0.1382
Epoch 57/100
9/9 [=====] - 0s 5ms/step - loss: 0.0894 - val_loss:
0.1441
Epoch 58/100
9/9 [=====] - 0s 5ms/step - loss: 0.0851 - val_loss:
0.1298
Epoch 59/100
9/9 [=====] - 0s 6ms/step - loss: 0.0798 - val_loss:
0.1414
Epoch 60/100
9/9 [=====] - 0s 6ms/step - loss: 0.0816 - val_loss:
0.1331
Epoch 61/100
9/9 [=====] - 0s 6ms/step - loss: 0.0864 - val_loss:
0.1366
Epoch 62/100
9/9 [=====] - 0s 6ms/step - loss: 0.0788 - val_loss:
0.1344
Epoch 63/100
9/9 [=====] - 0s 5ms/step - loss: 0.0743 - val_loss:
0.1407
Epoch 64/100
9/9 [=====] - 0s 5ms/step - loss: 0.0793 - val_loss:
0.1341
Epoch 65/100
9/9 [=====] - 0s 5ms/step - loss: 0.0773 - val_loss:
0.1359
Epoch 66/100
9/9 [=====] - 0s 5ms/step - loss: 0.0779 - val_loss:
0.1177
Epoch 67/100
9/9 [=====] - 0s 5ms/step - loss: 0.0757 - val_loss:
0.1357
Epoch 68/100
9/9 [=====] - 0s 5ms/step - loss: 0.0695 - val_loss:
0.1333
Epoch 69/100


9/9 [=====] - 0s 5ms/step - loss: 0.0783 - val_loss:
0.1351
Epoch 70/100
9/9 [=====] - 0s 5ms/step - loss: 0.0715 - val_loss:
0.1340
Epoch 71/100
9/9 [=====] - 0s 5ms/step - loss: 0.0739 - val_loss:
0.1339
Epoch 72/100
9/9 [=====] - 0s 5ms/step - loss: 0.0726 - val_loss:
0.1291
Epoch 73/100
9/9 [=====] - 0s 5ms/step - loss: 0.0819 - val_loss:
0.1262
Epoch 74/100
9/9 [=====] - 0s 5ms/step - loss: 0.0707 - val_loss:
0.1214
Epoch 75/100
9/9 [=====] - 0s 5ms/step - loss: 0.0782 - val_loss:
0.1518
Epoch 76/100
9/9 [=====] - 0s 5ms/step - loss: 0.0772 - val_loss:
0.1329
Epoch 77/100
9/9 [=====] - 0s 5ms/step - loss: 0.0732 - val_loss:
0.1285
Epoch 78/100
9/9 [=====] - 0s 5ms/step - loss: 0.0787 - val_loss:
0.1391
Epoch 79/100
9/9 [=====] - 0s 5ms/step - loss: 0.0753 - val_loss:
0.1365
Epoch 80/100
9/9 [=====] - 0s 5ms/step - loss: 0.0663 - val_loss:
0.1301
Epoch 81/100
9/9 [=====] - 0s 5ms/step - loss: 0.0667 - val_loss:
0.1288
Epoch 82/100
9/9 [=====] - 0s 5ms/step - loss: 0.0612 - val_loss:
0.1373
Epoch 83/100
9/9 [=====] - 0s 6ms/step - loss: 0.0617 - val_loss:
0.1409
Epoch 84/100
9/9 [=====] - 0s 5ms/step - loss: 0.0673 - val_loss:
0.1392
Epoch 85/100

9/9 [=====] - 0s 5ms/step - loss: 0.0690 - val_loss:
0.1255
Epoch 86/100
9/9 [=====] - 0s 5ms/step - loss: 0.0597 - val_loss:
0.1412
Epoch 87/100
9/9 [=====] - 0s 5ms/step - loss: 0.0688 - val_loss:
0.1393
Epoch 88/100
9/9 [=====] - 0s 5ms/step - loss: 0.0754 - val_loss:
0.1408
Epoch 89/100
9/9 [=====] - 0s 5ms/step - loss: 0.0673 - val_loss:
0.1413
Epoch 90/100
9/9 [=====] - 0s 6ms/step - loss: 0.0719 - val_loss:
0.1380
Epoch 91/100
9/9 [=====] - 0s 6ms/step - loss: 0.0717 - val_loss:
0.1393
Epoch 92/100
9/9 [=====] - 0s 5ms/step - loss: 0.0714 - val_loss:
0.1385
Epoch 93/100
9/9 [=====] - 0s 5ms/step - loss: 0.0658 - val_loss:
0.1469
Epoch 94/100
9/9 [=====] - 0s 5ms/step - loss: 0.0645 - val_loss:
0.1211
Epoch 95/100
9/9 [=====] - 0s 5ms/step - loss: 0.0691 - val_loss:
0.1591
Epoch 96/100
9/9 [=====] - 0s 5ms/step - loss: 0.0615 - val_loss:
0.1227
Epoch 97/100
9/9 [=====] - 0s 5ms/step - loss: 0.0656 - val_loss:
0.1408
Epoch 98/100
9/9 [=====] - 0s 5ms/step - loss: 0.0680 - val_loss:
0.1309
Epoch 99/100
9/9 [=====] - 0s 5ms/step - loss: 0.0564 - val_loss:
0.1417
Epoch 100/100
9/9 [=====] - 0s 5ms/step - loss: 0.0699 - val_loss:
0.1315

[66]: <keras.callbacks.History at 0x1e1897a5700>

```
[67]: stand_losses = pd.DataFrame(model_stand.history.history)
      stand_losses.plot()
```

[67]: <AxesSubplot:>



```
[68]: stand_predictions = model_stand.predict(X_test_stand)
```

```
[69]: mean_absolute_error(y_test_stand, stand_predictions)
```

[69]: 0.2528513299851795

```
[70]: np.sqrt(mean_squared_error(y_test_stand, stand_predictions))
```


[70]: 0.36256629216467884

```
[71]: explained_variance_score(y_test_stand, stand_predictions)
```

[71]: 0.8807082761768282

```
[72]: plt.scatter(y_test_stand, stand_predictions)
      plt.plot(y_test_stand, y_test_stand, 'r')
```

[72]: [<matplotlib.lines.Line2D at 0x1e1881c4760>]



```
[73]: # TODO: inverse scale predictions

temp_train = train.drop('SalePrice', axis=1)
temp_train['SalePrice'] = train['SalePrice']

scaler = StandardScaler()
scaler.fit(temp_train.drop(labels='Id', axis=1))
```

```
[73]: StandardScaler()
```

```
[74]: temp = X_test_stand.copy()
temp['Predictions'] = stand_predictions
```

```
[75]: temp_trans = pd.DataFrame(scaler.inverse_transform(temp), columns=temp.columns)
```

```
[76]: corrected_stand_pred = temp_trans['Predictions'].to_numpy()
```


```
[77]: temp_corrected_test = X_test_stand.copy()
temp_corrected_test['y_test'] = y_test_stand.to_numpy()
```

```
[78]: temp_trans_test = pd.DataFrame(scaler.inverse_transform(temp_corrected_test),
                                     columns=temp_corrected_test.columns)
```

```
[79]: corrected_stand_test = temp_trans_test['y_test'].to_numpy()
```

```
[80]: plt.scatter(corrected_stand_test, corrected_stand_pred)
plt.plot(corrected_stand_test, corrected_stand_test, 'r')
```

```
[80]: [<matplotlib.lines.Line2D at 0x1e18b0770a0>]
```



```
[81]: # inverse scaling working properly, make test predictions and apply inverse_
↳ transformation
```

```
[82]: scaled_predictions = model_stand.predict(test_standard)
```

```
[83]: temp = test_standard.copy()
temp['Predictions'] = scaled_predictions
```

```
[84]: temp_predictions = pd.DataFrame(scaler.inverse_transform(temp), columns=temp.
↳ columns)
```

```
[85]: predictions_array = temp_predictions['Predictions'].to_numpy()
```

```
[86]: predictions = pd.DataFrame({'Id': test.Id, 'SalePrice': predictions_array})
```

```
[87]: predictions.tail()
```

```
[87]:      Id      SalePrice
1454  2915  109484.807776
1455  2916   99962.820805
1456  2917  158730.260232
```

```
1457  2918  139104.945553
1458  2919  241892.581487
```

```
[88]: sample_submission = pd.read_csv('sample_submission.csv')
      sample_submission.shape
```

```
[88]: (1459, 2)
```

```
[89]: predictions.shape
```

```
[89]: (1459, 2)
```

```
[90]: predictions.to_csv('submission.csv', index=False)
```